

randomforest

November 20, 2024

```
[1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestRegressor
from datetime import datetime, timedelta
import seaborn as sns
```

0.0.1 Here we will provide a function for our Random Forest Regressor

```
[2]: def train_eval_rf():
    try:
        # Read data
        df = pd.read_csv('Superstore.csv', encoding='ISO-8859-1')

        # Create features
        features = ['Sales', 'Category', 'Sub-Category', 'Region', 'Segment']

        # Convert categorical variables to dummy variables
        df_encoded = pd.get_dummies(df[features], columns=['Category', 'Sub-Category', 'Region', 'Segment'])

        # Prepare X and y
        X = df_encoded
        y = df['Profit']

        # Split data
        X_train, X_test, y_train, y_test = train_test_split(
            X, y, test_size=0.2, random_state=42
        )

        # Create and train model
        rf_model = RandomForestRegressor(
            n_estimators=100,
```

```

        max_depth=None,
        min_samples_split=2,
        min_samples_leaf=1,
        random_state=42
    )

    rf_model.fit(X_train, y_train)

    # Make predictions
    y_pred_train = rf_model.predict(X_train)
    y_pred_test = rf_model.predict(X_test)

    # Calculate scores
    train_score = r2_score(y_train, y_pred_train)
    test_score = r2_score(y_test, y_pred_test)

    # Get feature importance
    feature_importance = pd.DataFrame({
        'feature': X_train.columns,
        'importance': rf_model.feature_importances_
    }).sort_values('importance', ascending=False)

    # Create feature importance plot
    plt.figure(figsize=(12, 6))
    plt.bar(feature_importance['feature'][:10],
    ↪feature_importance['importance'][:10])
    plt.xticks(rotation=45, ha='right')
    plt.title('Top 10 Most Important Features')
    plt.tight_layout()

    # Return the values
    return train_score, test_score, feature_importance

except Exception as e:
    print(f"An error occurred: {str(e)}")
    return None, None, None

# Run the analysis
train_score, test_score, feature_importance = train_eval_rf()

# Check if we got valid results
if train_score is not None:
    print(f"\nRandom Forest Results:")
    print(f"Training R2 score: {train_score:.4f}")
    print(f"Testing R2 score: {test_score:.4f}")
    print("\nTop 5 Most Important Features:")
    print(feature_importance.head())

```

```
plt.show()
else:
    print("Failed to run the analysis. Please check your data and try again.")
```


Random Forest Results:

Training R^2 score: 0.9120

Testing R^2 score: -0.6703

Top 5 Most Important Features:

| | feature | importance |
|----|----------------------|------------|
| 0 | Sales | 0.831803 |
| 7 | Sub-Category_Binders | 0.026070 |
| 21 | Region_Central | 0.018102 |
| 26 | Segment_Corporate | 0.014203 |
| 25 | Segment_Consumer | 0.012612 |



0.0.2 We are going to tweak our model a bit and see if we can get a better mix of training and test results.

```
[ ]: from sklearn.model_selection import train_test_split
     from sklearn.metrics import r2_score
     from sklearn.metrics import roc_curve, auc
```

```
[4]: def train_eval_rf_models():
     # Read and prepare data
     df = pd.read_csv('Superstore.csv', encoding='ISO-8859-1')

     # Feature engineering
```

```

df['Profit_Margin'] = df['Profit'] / df['Sales']
df['Discount_Amount'] = df['Sales'] * df['Discount']

features = [
    'Sales', 'Discount', 'Quantity',
    'Profit_Margin', 'Discount_Amount',
    'Category', 'Sub-Category', 'Region',
    'Segment', 'Ship Mode'
]

# Prepare data
df_encoded = pd.get_dummies(df[features],
                             columns=['Category', 'Sub-Category',
                                      'Region', 'Segment', 'Ship Mode'])

X = df_encoded
y = df['Profit']

# Split data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Define models
rf_model_base = RandomForestRegressor(
    max_depth=None,
    max_features="sqrt",
    max_leaf_nodes=None,
    min_samples_leaf=1,
    min_samples_split=2,
    min_weight_fraction_leaf=0.0,
    n_estimators=8,
    n_jobs=1,
    oob_score=False,
    random_state=None,
    verbose=0,
    warm_start=False
)

rf_model_V1 = RandomForestRegressor(
    n_estimators=25,
    max_features="sqrt",
    bootstrap=True,
    oob_score=True,
    random_state=42
)

rf_model_V2 = RandomForestRegressor(

```

```

        n_estimators=15,
        max_features="sqrt",
        min_samples_leaf=2,
        min_samples_split=3,
        random_state=42
    )

rf_model_V3 = RandomForestRegressor(
    n_estimators=20,
    max_features="sqrt",
    max_leaf_nodes=None,
    min_samples_leaf=2,
    bootstrap=True,
    random_state=42
)

# Dictionary to store results
results = {}

# Train and evaluate each model
for name, model in [
    ('Base', rf_model_base),
    ('V1', rf_model_V1),
    ('V2', rf_model_V2),
    ('V3', rf_model_V3)
]:
    # Train model
    model.fit(X_train, y_train)

    # Make predictions
    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)

    # Calculate scores
    train_score = r2_score(y_train, y_pred_train)
    test_score = r2_score(y_test, y_pred_test)

    # Store feature importance
    feature_importance = pd.DataFrame({
        'feature': X_train.columns,
        'importance': model.feature_importances_
    }).sort_values('importance', ascending=False)

    results[name] = {
        'train_score': train_score,

```

```

        'test_score': test_score,
        'feature_importance': feature_importance
    }

    return results

# Run all models and get results
results = train_eval_rf_models()

# Print results
for model_name, model_results in results.items():
    print(f"\nModel {model_name} Results:")
    print(f"Training R2 score: {model_results['train_score']:.4f}")
    print(f"Testing R2 score: {model_results['test_score']:.4f}")
    print("\nTop 5 Most Important Features:")
    print(model_results['feature_importance'].head())
    print("-" * 50)

# Visualize comparison
plt.figure(figsize=(12, 6))
x = np.arange(len(results))
width = 0.35

plt.bar(x - width/2, [r['train_score'] for r in results.values()], width,
        label='Training Score')
plt.bar(x + width/2, [r['test_score'] for r in results.values()], width,
        label='Testing Score')

plt.xlabel('Model Version')
plt.ylabel('R2 Score')
plt.title('Model Performance Comparison')
plt.xticks(x, results.keys())
plt.legend()
plt.tight_layout()
plt.show()

```

Model Base Results:

Training R² score: 0.9479

Testing R² score: 0.6495

Top 5 Most Important Features:

| | feature | importance |
|----|----------------------|------------|
| 0 | Sales | 0.407304 |
| 4 | Discount_Amount | 0.153583 |
| 3 | Profit_Margin | 0.114142 |
| 14 | Sub-Category_Copiers | 0.072162 |

| | | |
|---|----------|----------|
| 2 | Quantity | 0.066964 |
|---|----------|----------|

Model V1 Results:

Training R^2 score: 0.9655

Testing R^2 score: 0.6230

Top 5 Most Important Features:

| | feature | importance |
|----|----------------------|------------|
| 0 | Sales | 0.502322 |
| 4 | Discount_Amount | 0.121559 |
| 3 | Profit_Margin | 0.115381 |
| 14 | Sub-Category_Copiers | 0.047699 |
| 2 | Quantity | 0.040134 |

Model V2 Results:

Training R^2 score: 0.8949

Testing R^2 score: 0.7151

Top 5 Most Important Features:

| | feature | importance |
|----|----------------------|------------|
| 0 | Sales | 0.514651 |
| 4 | Discount_Amount | 0.168590 |
| 3 | Profit_Margin | 0.088355 |
| 14 | Sub-Category_Copiers | 0.066833 |
| 2 | Quantity | 0.031925 |

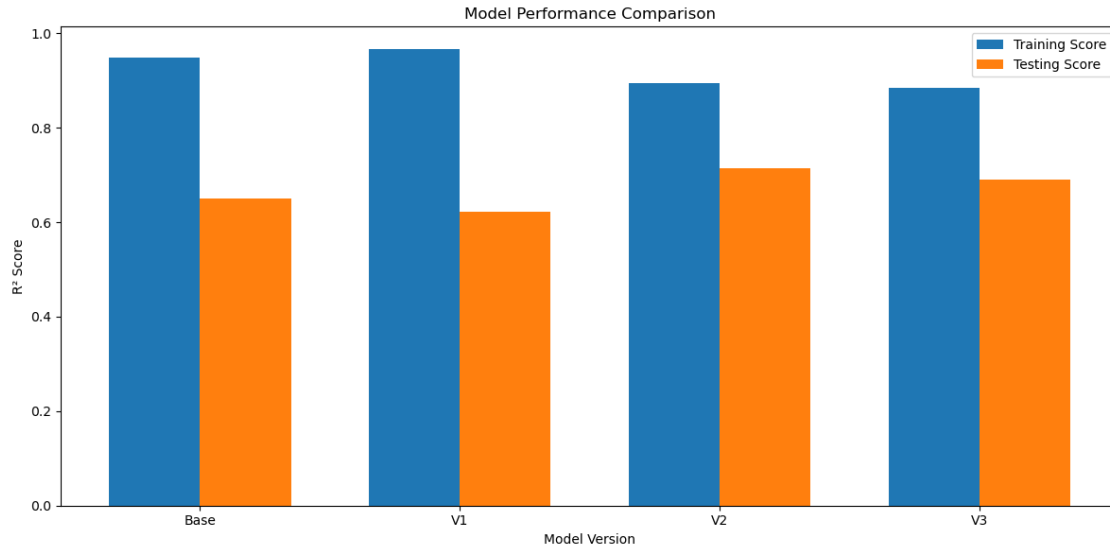
Model V3 Results:

Training R^2 score: 0.8844

Testing R^2 score: 0.6893

Top 5 Most Important Features:

| | feature | importance |
|----|----------------------|------------|
| 0 | Sales | 0.487046 |
| 4 | Discount_Amount | 0.152881 |
| 3 | Profit_Margin | 0.101438 |
| 14 | Sub-Category_Copiers | 0.069265 |
| 1 | Discount | 0.033524 |



```
[5]: def train_eval_rf_stable():
    # Set random seed for entire process
    np.random.seed(42)

    # Read and prepare data
    df = pd.read_csv('Superstore.csv', encoding='ISO-8859-1')

    # Feature engineering
    df['Profit_Margin'] = df['Profit'] / df['Sales']
    df['Discount_Amount'] = df['Sales'] * df['Discount']

    features = [
        'Sales', 'Discount', 'Quantity',
        'Profit_Margin', 'Discount_Amount',
        'Category', 'Sub-Category', 'Region',
        'Segment', 'Ship Mode'
    ]

    # Prepare data
    df_encoded = pd.get_dummies(df[features],
                                columns=['Category', 'Sub-Category',
                                           'Region', 'Segment', 'Ship Mode'])

    X = df_encoded
    y = df['Profit']

    # Split data with fixed random state
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42)
```



```

)

# Define model with fixed random state
rf_model = RandomForestRegressor(
    max_depth=None,
    max_features="sqrt",
    max_leaf_nodes=None,
    min_samples_leaf=1,
    min_samples_split=2,
    min_weight_fraction_leaf=0.0,
    n_estimators=10,
    n_jobs=1,
    random_state=42, # Fixed random state
    bootstrap=True # Explicitly set bootstrap
)

# Train model
rf_model.fit(X_train, y_train)

# Make predictions
y_pred_train = rf_model.predict(X_train)
y_pred_test = rf_model.predict(X_test)

# Calculate scores
train_score = r2_score(y_train, y_pred_train)
test_score = r2_score(y_test, y_pred_test)

# Get feature importance
feature_importance = pd.DataFrame({
    'feature': X_train.columns,
    'importance': rf_model.feature_importances_
}).sort_values('importance', ascending=False)

return train_score, test_score, feature_importance

# Run the model multiple times to verify stability
print("Running model multiple times to verify stability:\n")

for i in range(3):
    train_score, test_score, feature_importance = train_eval_rf_stable()
    print(f"Run {i+1}:")
    print(f"Training R2 score: {train_score:.4f}")
    print(f"Testing R2 score: {test_score:.4f}")
    print(f"\nTop 5 Most Important Features:")
    print(feature_importance.head())
    print("-" * 50 + "\n")

```

Running model multiple times to verify stability:

Run 1:

Training R^2 score: 0.9756

Testing R^2 score: 0.6164

Top 5 Most Important Features:

| | feature | importance |
|----|----------------------|------------|
| 0 | Sales | 0.498980 |
| 4 | Discount_Amount | 0.128803 |
| 3 | Profit_Margin | 0.100625 |
| 14 | Sub-Category_Copiers | 0.051662 |
| 2 | Quantity | 0.038025 |

Run 2:

Training R^2 score: 0.9756

Testing R^2 score: 0.6164

Top 5 Most Important Features:

| | feature | importance |
|----|----------------------|------------|
| 0 | Sales | 0.498980 |
| 4 | Discount_Amount | 0.128803 |
| 3 | Profit_Margin | 0.100625 |
| 14 | Sub-Category_Copiers | 0.051662 |
| 2 | Quantity | 0.038025 |

Run 3:

Training R^2 score: 0.9756

Testing R^2 score: 0.6164

Top 5 Most Important Features:

| | feature | importance |
|----|----------------------|------------|
| 0 | Sales | 0.498980 |
| 4 | Discount_Amount | 0.128803 |
| 3 | Profit_Margin | 0.100625 |
| 14 | Sub-Category_Copiers | 0.051662 |
| 2 | Quantity | 0.038025 |

```
[6]: def test_n_estimators(n_range):  
      # Set random seed  
      np.random.seed(42)  
  
      # Read and prepare data  
      df = pd.read_csv('Superstore.csv', encoding='ISO-8859-1')
```

```

# Feature engineering
df['Profit_Margin'] = df['Profit'] / df['Sales']
df['Discount_Amount'] = df['Sales'] * df['Discount']

features = [
    'Sales', 'Discount', 'Quantity',
    'Profit_Margin', 'Discount_Amount',
    'Category', 'Sub-Category', 'Region',
    'Segment', 'Ship Mode'
]

# Prepare data
df_encoded = pd.get_dummies(df[features],
                             columns=['Category', 'Sub-Category',
                                       'Region', 'Segment', 'Ship Mode'])

X = df_encoded
y = df['Profit']

# Split data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

results = []

# Test different numbers of trees
for n in n_range:
    rf_model = RandomForestRegressor(
        n_estimators=n,
        max_depth=None,
        max_features="sqrt",
        max_leaf_nodes=None,
        min_samples_leaf=1,
        min_samples_split=2,
        random_state=42
    )

    rf_model.fit(X_train, y_train)

    train_score = r2_score(y_train, rf_model.predict(X_train))
    test_score = r2_score(y_test, rf_model.predict(X_test))

    results.append({
        'n_estimators': n,
        'train_score': train_score,
        'test_score': test_score
    })

```

```

    })

    return pd.DataFrame(results)

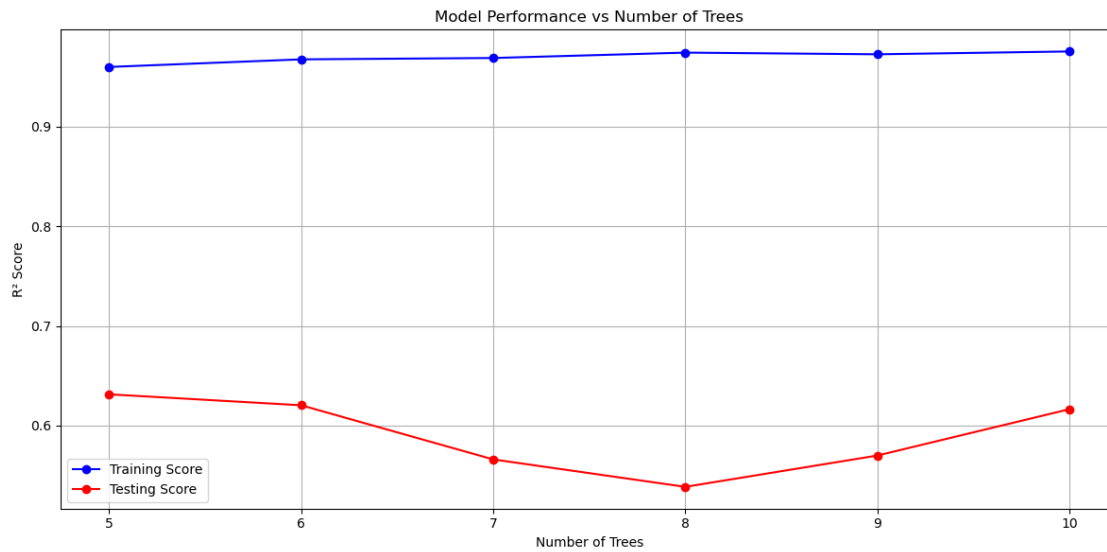
# Test range of trees around 7
n_range = [5, 6, 7, 8, 9, 10]
results_df = test_n_estimators(n_range)

# Plot results
plt.figure(figsize=(12, 6))
plt.plot(results_df['n_estimators'], results_df['train_score'],
         'b-o', label='Training Score')
plt.plot(results_df['n_estimators'], results_df['test_score'],
         'r-o', label='Testing Score')
plt.xlabel('Number of Trees')
plt.ylabel('R2 Score')
plt.title('Model Performance vs Number of Trees')
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.show()

# Print detailed results
print("\nDetailed Results:")
print(results_df.round(4))

# Find best number of trees
best_test = results_df.loc[results_df['test_score'].idxmax()]
print(f"\nBest Performance:")
print(f"Number of Trees: {best_test['n_estimators']}")
print(f"Test Score: {best_test['test_score']:.4f}")
print(f"Train Score: {best_test['train_score']:.4f}")

```



Detailed Results:

| | n_estimators | train_score | test_score |
|---|--------------|-------------|------------|
| 0 | 5 | 0.9600 | 0.6313 |
| 1 | 6 | 0.9677 | 0.6203 |
| 2 | 7 | 0.9691 | 0.5659 |
| 3 | 8 | 0.9745 | 0.5383 |
| 4 | 9 | 0.9728 | 0.5698 |
| 5 | 10 | 0.9756 | 0.6164 |

Best Performance:

Number of Trees: 5.0

Test Score: 0.6313

Train Score: 0.9600

[]: